



AI CONTENT CREATION SYSTEM

UXAI Content Creator

Brand-aligned content from real source material, not generic AI output.

by
Nelly, Gretel & Kimia

July 2026

INTRODUCTION

Generic AI content is easy to distrust.



Generic AI output

Anyone can prompt it.



Grounded in real sources

A real book · a real framework · real data.



Two markdown knowledge bases + non-RAG context injection

KNOWLEDGE BASE

Two knowledge bases. Six files.

PRIMARY — BRAND



book_framework.md



ai_interaction_map
.md



bio_credentials.md



past_content_samples
.md

SECONDARY — RESEARCH



industry_trends.md



competitor_positioning.md

KEYWORD SCORING -> NO VECTOR STORE



Load

all six files



Score

keyword overlap



Select

1–2 primary + 1 secondary



Inject

straight into the prompt

No RAG. On purpose.



Six files fit in one prompt

no scale problem to solve



Narrow, predictable queries

retrieval adds cost, not value




Two-day scope

time went into uniqueness instead

Revisit past 15–20 documents. `Full defense: rag_decision.md`

LIVE DEMO

The pipeline, end to end.



Nothing auto-publishes. Every decision is logged with a timestamp.

UNIQUENESS

Four strategies. Real evidence.



Human in the Loop

Logged review checkpoint



Style Variation

3 genuinely different templates



Contextual Placement

Real names. Real numbers.




Iterative Prompting

Before → fix → after

PROOF

We caught our own fabrication.



~~“Clients report clearer understanding...”~~

invented — stated as fact



prompt refined: hypothetical framing required



“While specific metrics aren't available here...”

confirmed on the next generation

SAME TOPIC · SAME FORMAT · NO CONTEXT FOR CHATGPT



Our system

Named the framework. Cited real data.



ChatGPT

A listicle anyone could write.

Same topic, same format: Chatbot UX

Our system

LinkedIn Post · chatbot ux

Most chatbot UX flaws trace back to treating the interface as the whole system — a surface-level design problem. But AI interaction is a stack of distinct layers, each with its own risks and breakdowns.

The “interface is not the system” is more than a catchphrase. It’s a framework to diagnose why chatbots fail adoption despite being shipped widely. Gartner reports 80% of enterprises now embed AI agents, yet only 31% see real production usage. Often, poor chatbot UX reflects failure to manage what happens beneath the surface: perception errors, opaque reasoning, incomplete memory, or missing controls for humans to intervene.

A good chatbot UX doesn’t just enable conversation; it clarifies what the AI does and why, surfaces confidence cues, and integrates oversight points. Transparency isn’t an add-on; it’s a governance decision intertwined with every AI layer.

Next time you design or evaluate a chatbot, ask: Which AI layers are exposed or hidden? Where does trust break down? Closing these gaps drives better adoption and real value — not just prettier chat windows.

How deeply does your chatbot’s UX reach beneath the surface?

ChatGPT

Redacción

Good chatbot UX isn’t about making a bot sound human. It’s about helping people reach their goal with as little effort as possible.

The best chatbot experiences are usually the ones users barely notice:

- Clear, simple language
- Fast and relevant responses
- Smooth handoff to a human when needed
- Context remembered throughout the conversation
- Transparency about what the bot can and can’t do

As AI capabilities continue to evolve, great UX is becoming an even bigger differentiator. A powerful model alone isn’t enough—how people experience that intelligence is what builds trust.




The question is no longer *“Can we automate this?”* but *“How can we make this interaction genuinely useful?”*

What’s the best chatbot experience you’ve had recently?

#ChatbotUX #ConversationalAI #CustomerExperience #UXDesign #ArtificialIntelligence #CX

PROCESS

How we worked: the board.



Small, modular, inspectable.



Python 3 + OpenAI SDK



Markdown KB - no vector store



Codex fenced off: scaffolding only

`document_processor.py`



`knowledge_base.py`



`prompt_templates.py`



`llm_integration.py`

`content_pipeline.py`



`review.py`



`main.py / app.py`

• UXAI Content Creator

HistorySTEP 1 OF 4

NEW DRAFT

What are we writing about?

Enter your content topic

AI-Driven Product Design

Continue



• UXAI Content Creator

HistorySTEP 2 OF 4

Topic: AI-Driven Product Design · edit

Choose a format

Thought Leadership

Long-form argument with a clear thesis, for deep, considered takes.


LinkedIn Post

Short, hook-driven, mobile-first. One idea, one takeaway.

Case Study Summary

Structured Problem / Approach / Insight / Impact format.

Generate



LinkedIn Post · AI driven product design

Most AI-driven product designs start by treating “AI” as a single magic box. That’s the problem.

AI isn’t one thing. It’s a stack of capabilities—16 layers grouped across Perception, Generation, Communication, Reasoning, Memory, and Action. Each layer works differently, fails differently, and requires distinct oversight.

Ignoring this complexity leads to products where users guess what AI did—or worse, if it even did anything at all. Oversight and transparency aren’t optional extras; they’re design imperatives embedded deep into the stack.

Designers must extend their role beyond crafting interfaces. Understanding where control breaks down, how reasoning surfaces, and where trust erodes becomes vital.

Products like Spotify or Tesla don’t just “use AI.” They orchestrate layers to balance automation with human authority. That’s the real design challenge.

How are you mapping AI’s layers in your product design—not just slapping on an AI feature? #AIInteractionMap #ProductDesign #TrustInAI

QUALITY CHECK

Passed

Approve

Mark as final

Request edits

Send back with notes

Reject

Discard this draft

LinkedIn Post · AI driven product design

Most AI-driven product designs start by treating "AI" as a single magic box. That's the problem.

AI isn't one thing. It's a stack of capabilities—16 layers grouped across Perception, Generation, Communication, Reasoning, Memory, and Action. Each layer works differently, fails differently, and requires distinct oversight.

Ignoring this complexity leads to products where users guess what AI did—or worse, if it even did anything at all. Oversight and transparency aren't optional extras; they're design imperatives embedded deep into the stack.

Designers must extend their role beyond crafting interfaces. Understanding where control breaks down, how reasoning surfaces, and where trust erodes becomes vital.

Products like Spotify or Tesla don't just "use AI." They orchestrate layers to balance automation with human authority. That's the real design challenge.

How are you mapping AI's layers in your product design—not just slapping on an AI feature? #AIInteractionMap #ProductDesign #TrustInAI

QUALITY CHECK

Passed

Approve

Mark as final

Request edits

Send back with notes


Reject

Discard this draft

Describe the changes you'd like

add real statistics to the post

Resubmit



Revising your LinkedIn Post...

Rewriting the draft...



LinkedIn Post · AI driven product design

REVISED

* add real statistics to the post*

Most AI-driven product designs start by treating "AI" as a single magic box. That's the problem.

AI isn't one thing. It's a stack of capabilities—16 layers grouped across Perception, Generation, Communication, Reasoning, Memory, and Action. Each layer works differently, fails differently, and requires distinct oversight.

Gartner data shows 80% of enterprise applications shipped in Q1 2026 embed at least one AI agent, but only 31% see actual production usage. This gap reveals how ignoring AI's complexity leads to products where users guess what AI did—or worse, if it even did anything at all. Oversight and transparency aren't optional extras; they're design imperatives embedded deep into the stack.

Designers must extend their role beyond crafting interfaces. Understanding where control breaks down, how reasoning surfaces, and where trust erodes becomes vital.

Products like Spotify or Tesla don't just "use AI." They orchestrate layers to balance automation with human authority. That's the real design challenge.

How are you mapping AI's layers in your product design—not just slapping on an AI feature? #AIInteractionMap #ProductDesign #TrustInAI

QUALITY CHECK

Review Recommended

Issues

- The post includes a percentage or statistic.

Suggestions

- Check that any statistic is true and has a reliable source.

Approve


Mark as final

Request edits

Send back with notes

Reject

Discard this draft



Verified statistics across multiple sources and distinguished related concepts

Good news: the stat checks out. Multiple independent sources (Gartner directly, plus corroborating pickups citing Gartner and S&P Global Market Intelligence) confirm the same two figures: 80% of enterprise applications shipped or updated in Q1 2026 embed at least one AI agent, per Gartner, up from 33% in 2024, and 31% of enterprises have at least one AI agent in production, per S&P Global Market Intelligence and McKinsey.

Your post's phrasing matches this accurately, "80% of enterprise applications shipped in Q1 2026 embed at least one AI agent, but only 31% see actual production usage" is a fair, correct paraphrase of the real numbers.

One thing worth being careful about, since it came up in these search results: there are several different Gartner stats floating around with similar-looking numbers, easy to accidentally mix up:

- The 80%/31% gap you used (real, Q1 2026, correctly sourced)
- A separate 2023 Gartner prediction that "80% of enterprises will have used GenAI APIs by 2026" (different metric entirely, about API usage, not shipped agents)
- A separate Gartner prediction that "40% of enterprise apps will feature task-specific AI agents by end of 2026" (a forward-looking prediction, not the Q1 2026 measured 80%)

Your post correctly uses the first one and doesn't conflate it with the others, but worth double-checking any future posts against `industry_trends.md` directly rather than from memory, since these numbers are genuinely easy to mix up even when writing carefully.

Fact Check

APPROVED

Content approved

This draft is marked final and logged to the review history.

Most AI-driven product designs start by treating "AI" as a single magic box. That's the problem.

AI isn't one thing. It's a stack of capabilities—16 layers grouped across Perception, Generation, Communication, Reasoning, Memory, and Action. Each layer works differently, fails differently, and requires distinct oversight.

Gartner data shows 80% of enterprise applications shipped in Q1 2026 embed at least one AI agent, but only 31% see actual production usage. This gap reveals how ignoring AI's complexity leads to products where users guess what AI did—or worse, if it even did anything at all. Oversight and transparency aren't optional extras; they're design imperatives embedded deep into the stack.

Designers must extend their role beyond crafting interfaces. Understanding where control breaks down, how reasoning surfaces, and where trust erodes becomes vital.

Products like Spotify or Tesla don't just "use AI." They orchestrate layers to balance automation with human authority. That's the real design challenge.

How are you mapping AI's layers in your product design—not just slapping on an AI feature? #AIInteractionMap #ProductDesign #TrustInAI

Start a new draft

Download

Finish

What actually went wrong.



Path resolution

Silent CWD bug → anchored paths



Python 3.9

str | None → Optional[str]



Our own fabrication

Caught it. Fixed it. Proved it.

THE CORE FINDING

The difference between generic AI content and brand-aligned content isn't a better model. It's better context, and a Human In The Loop that catches mistakes before they ship instead of after!

Questions?

github.com/np8350-prog/uxai-content-creator